In [30]:
```
import pandas as pd
import numpy as np
```

In [31]:
```
df = pd.read_csv('E:/Work/P/Workspace/SupL_1/import2.csv',names =['symboling',
 'normalized-losses', 'make', 'fuel-type', 'aspiration', 'num-of-doors', 'body
-style',
          'drive-wheels', 'engine-location', 'wheel-base', 'length', 'width',
 'height', 'curb-weight',
          'engine-type', 'num-of-cylinders', 'engine-size', 'fuel-system', 'b
ore', 'stroke',
          'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg', 'highway
-mpg', 'price'])
df_head=df.head(10)
```

In [33]:
```
df_head
```

Out[33]:

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location
0	3	NaN	alfa-romero	gas	std	two	convertible	rwd	front
1	3	NaN	alfa-romero	gas	std	two	convertible	rwd	front
2	1	NaN	alfa-romero	gas	std	two	hatchback	rwd	front
3	2	164.0	audi	gas	std	four	sedan	fwd	front
4	2	164.0	audi	gas	std	four	sedan	4wd	front
5	2	NaN	audi	gas	std	two	sedan	fwd	front
6	1	158.0	audi	gas	std	four	sedan	fwd	front
7	1	NaN	audi	gas	std	four	wagon	fwd	front
8	1	158.0	audi	gas	turbo	four	sedan	fwd	front
9	0	NaN	audi	gas	turbo	two	hatchback	4wd	front

10 rows × 26 columns

In [34]:
```
#Q 2:Explain the problem statement. What are you predicting and what attribute
s you have to predict?
#Ans : We can use this data to predict the price of the car based on various a
ttributes
```

In [35]:
```
#Q 3:Browse a sample record from the dataframe. Are there any missing values?
df.iloc[14]
#The normalized losses value is missing for this tuple
```

Out[35]:
```
symboling               1
normalized-losses     NaN
make                  bmw
fuel-type             gas
aspiration            std
num-of-doors         four
body-style          sedan
drive-wheels          rwd
engine-location     front
wheel-base          103.5
length                189
width                66.9
height               55.7
curb-weight          3055
engine-type           ohc
num-of-cylinders      six
engine-size           164
fuel-system          mpfi
bore                 3.31
stroke               3.19
compression-ratio       9
horsepower            121
peak-rpm             4250
city-mpg               20
highway-mpg            25
price               24565
Name: 14, dtype: object
```

In [36]:
```
# #Q 4: How many records are available in the data set and how many attributes.
# Do you think the depth (number of records) is sufficient given the breadth?
# In other words, is the sample likely to be a good representative of the universe?
#Ans : Number of Records - 205
df.shape[0]
```

Out[36]:
```
205
```

In [37]:
```
#Ans : Number of Columns - 26
df.shape[1]
```

Out[37]:
```
26
```

In [ ]:
```
#Que: Do you think the depth (number of records) is sufficient given the bread
th?
#In other words, is the sample likely to be a good representative of the unive
rse?

#Ans : No the number of records is not sufficient for precise prediction compa
red to
#the given number of attributes. The sample is a mediocre representative of th
e universe.
```

In [ ]:
```
#Que : Analyse the data distribution for the various attributes and share your
 observations ?
```

In [53]:
```
df['price'].corr(df['symboling'])
```

Out[53]: -0.082391187169623584

In [54]:
```
df['price'].corr(df['wheel-base'])
```

Out[54]: 0.58464182226550787

In [55]:
```
df['price'].corr(df['length' ])
```

Out[55]: 0.69062838044836405

In [56]:
```
df['price'].corr(df['width'])
```

Out[56]: 0.75126534405226697

In [57]:
```
df['price'].corr(df['height'])
```

Out[57]: 0.13548630756805977

In [58]:
```
df['price'].corr(df['curb-weight'])
```

Out[58]: 0.83441452577028474

In [59]:
```
df['price'].corr(df['engine-size'])
```

Out[59]: 0.87233516744551831

In [60]:
```
df['price'].corr(df['bore'])
```

Out[60]: 0.54343586641885455

In [61]:
```
df['price'].corr(df['stroke'])
```

Out[61]: 0.082309827389704937

In [62]:
```
df['price'].corr(df['compression-ratio'])
```

Out[62]: 0.07110732668194146

In [63]:
```
df['price'].corr(df['horsepower'])
```

Out[63]: 0.81053308213220654

In [64]:
```
df['price'].corr(df['highway-mpg'])
```

Out[64]: -0.70469226505895299

In [65]:
```
#Ans : The following attributes have positive correlation with attribute prici
ng.
#(i.e the values that impact the price of the car)
        #wheel-base
        #length
        #width
        #height
        #curb-weight
        #engine-size
        #bore
        #stroke
        #compression-ratio
        #horsepower
#The following attributes have negative correlation with attribute pricing.
#(i.e the values that don't impact the price of the car)
        #symboling
        #highway-mpg
```

In [66]:
```
#Que : Are there any independent attributes which have |R| close to 1?
#Ans : engine-size attribute has |R| close to 1
```

In [67]:
```
#Que : Which attributes seem to have stronger relation with the dependent vari
able (Price of the car)?
#Ans : horsepower,engine-size and curb-weight
```

In [ ]:
```
#Que : Given the above analysis, which algorithm is likely to give a better ac
curacy? Why?
#Ans : logistic regression
```